

Loading data and preprocessing

In []:

```
!pip install lime
```

```
Requirement already satisfied: lime in /usr/local/lib/python3.7/dist-packages (0.2.0.1)
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from lime) (4.62.3)
Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages (from lime) (1.4.1)
Requirement already satisfied: scikit-image>=0.12 in /usr/local/lib/python3.7/dist-packages (from lime) (0.18.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.7/dist-packages (from lime) (3.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from lime) (1.21.5)
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.7/dist-packages (from lime) (1.0.2)
Requirement already satisfied: imageio>=2.3.0 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (2.4.1)
Requirement already satisfied: pillow!=7.1.0,!>=7.1.1,>=4.3.0 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (7.1.2)
Requirement already satisfied: PyWavelets>=1.1.1 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (1.2.0)
Requirement already satisfied: tifffile>=2019.7.26 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (2021.11.2)
Requirement already satisfied: networkx>=2.0 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (2.6.3)
Requirement already satisfied: pyparsing!=2.0.4,!>=2.1.2,!>=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (3.0.7)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (2.8.2)
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (0.11.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (1.3.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1->matplotlib->lime) (1.15.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.18->lime) (1.1.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.18->lime) (3.1.0)
```

In []:

```
!pip install tensorflow_text
```

```
Requirement already satisfied: tensorflow_text in /usr/local/lib/python3.7/dist-packages (2.8.1)
Requirement already satisfied: tensorflow<2.9,>=2.8.0 in /usr/local/lib/python3.7/dist-packages (from tensorflow_text) (2.8.0)
Requirement already satisfied: tensorflow-hub>=0.8.0 in /usr/local/lib/python3.7/dist-packages (from tensorflow_text) (0.12.0)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (3.10.0.2)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (1.13.3)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (0.24.0)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (3.3.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (1.43.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (1.1.0)
Requirement already satisfied: tensorboard<2.9,>=2.8 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (2.8.0)
Requirement already satisfied: tf-estimator-nightly==2.8.0.dev20211122109 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (2.8.0.dev20211122109)
Requirement already satisfied: protobuf>=3.9.2 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (3.17.3)
Requirement already satisfied: gast>=0.2.1 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (0.5.3)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (0.2.0)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (1.21.5)
Requirement already satisfied: setuptools in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (57.4.0)
Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (3.10.0)
```

```

8.0->tensorflow_text) (3.1.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (1.6.3)
Requirement already satisfied: keras<2.9,>=2.8.0rc0 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (2.8.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (1.15.0)
Requirement already satisfied: keras-preprocessing>=1.1.1 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (1.1.2)
Requirement already satisfied: absl-py>=0.4.0 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (1.0.0)
Requirement already satisfied: libclang>=9.0.1 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (13.0.0)
Requirement already satisfied: flatbuffers>=1.12 in /usr/local/lib/python3.7/dist-packages (from tensorflow<2.9,>=2.8.0->tensorflow_text) (2.0)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.7/dist-packages (from astunparse>=1.6.0->tensorflow<2.9,>=2.8.0->tensorflow_text) (0.37.1)
Requirement already satisfied: cached-property in /usr/local/lib/python3.7/dist-packages (from h5py>=2.9.0->tensorflow<2.9,>=2.8.0->tensorflow_text) (1.5.2)
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.7/dist-packages (from tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (1.0.1)
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (0.6.1)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.7/dist-packages (from tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (1.35.0)
Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (1.8.1)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.7/dist-packages (from tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (3.3.6)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.7/dist-packages (from tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (2.23.0)
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/local/lib/python3.7/dist-packages (from tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (0.4.6)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (0.2.8)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (4.8)
Requirement already satisfied: cachetools<5.0,>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (4.2.4)
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.7/dist-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (1.3.1)
Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.7/dist-packages (from markdown>=2.6.8->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (4.11.0)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from importlib-metadata>=4.4->markdown>=2.6.8->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (3.7.0)
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /usr/local/lib/python3.7/dist-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (0.4.8)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (2021.10.8)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (1.24.3)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (3.0.4)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (2.10)
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.7/dist-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorboard<2.9,>=2.8->tensorflow<2.9,>=2.8.0->tensorflow_text) (3.2.0)

```

In []:

```

import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from time import time
warnings.filterwarnings("ignore")
%matplotlib inline
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import re
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')

```

```

from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Flatten, Embedding, Input, Dropout, LSTM
from keras.utils.np_utils import to_categorical
from tensorflow.python.keras.callbacks import TensorBoard
from sklearn.model_selection import train_test_split
import tensorflow as tf
import tensorflow_hub as hub
import tensorflow_text as text
import pandas as pd
from sklearn.metrics import f1_score, confusion_matrix
from keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelCheckpoint
from lime.lime_text import LimeTextExplainer
from keras.models import load_model

```

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!

```

In []:

```

from google.colab import drive

drive.mount('/content/drive')

```

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```

In []:

```
%cd /content/drive/MyDrive/kaggle_toxic/
```

```
/content/drive/MyDrive/kaggle_toxic
```

In []:

```
all_data = pd.read_csv('all_data.csv')
```

In []:

```
all_data.head(5)
```

Out[]:

	id	comment_text	split	created_date	publication_id	parent_id	article_id	rating	funny	wow	sad	likes	disagree	toxicity	sev
0	1083994	He got his money... now he lies in wait till a...	train	2017-03-06 15:21:53.675241+00	21	NaN	317120	approved	0	0	0	2	0	0.373134	
1	650904	Mad dog will surely put the liberals in mental...	train	2016-12-02 16:44:21.329535+00	21	NaN	154086	approved	0	0	1	2	0	0.605263	
2	5902188	And Trump continues his lifelong cowardice by ...	train	2017-09-05 19:05:32.341360+00	55	NaN	374342	approved	1	0	2	3	7	0.666667	
3	7084460	"while arresting a man for resisting arrest".\...	test	2016-11-01 16:53:33.561631+00	13	NaN	149218	approved	0	0	0	0	0	0.815789	
4	5410943	Tucker and Paul are both total bad ass mofo's.	train	2017-06-14 05:08:21.997315+00	21	NaN	344096	approved	0	0	0	1	0	0.550000	

◀ | ▶

In []:

```

toxic = []
#making comments which have probability more than 0.5 as toxic and marking them as 1 while non-toxic as 0
for i in all_data['toxicity']:
    if i > 0.5:
        toxic.append(1)

```

```

else:
    toxic.append(0)

all_data['toxic_binary'] = toxic

```

In []:

```

all_data['sub_toxic'] = all_data[['severe_toxicity','obscene','sexual_explicit', 'identity_attack','insult',

In []:

```

```

sub_toxic = []
for j in range(len(all_data)):
    if all_data['toxic_binary'][j] == 1:
        if all_data['sub_toxic'][j] == 'severe_toxicity':
            sub_toxic.append(6)
        if all_data['sub_toxic'][j] == 'obscene':
            sub_toxic.append(5)
        if all_data['sub_toxic'][j] == 'sexual_explicit':
            sub_toxic.append(4)
        if all_data['sub_toxic'][j] == 'identity_attack':
            sub_toxic.append(3)
        if all_data['sub_toxic'][j] == 'insult':
            sub_toxic.append(2)
        if all_data['sub_toxic'][j] == 'threat':
            sub_toxic.append(1)
    if all_data['toxic_binary'][j] == 0:
        sub_toxic.append(0)

```

```

all_data['sub_toxic'] = sub_toxic

```

In []:

```

stop = set(stopwords.words('english'))

def clean(text):
    text_token = word_tokenize(text)
    filtered_text = ' '.join([w.lower() for w in text_token if w.lower() not in stop and len(w) > 2])
    filtered_text = filtered_text.replace(r"^[a-zA-Z]+", '')
    text_only = re.sub(r'\b\d+\b', '', filtered_text)
    clean_text = text_only.replace(',', '').replace('.', '').replace(':', '')
    return clean_text

```

In []:

```

all_data['clean_comment'] = [clean(str(x)) for x in all_data['comment_text']]

```

Splitting Data

In []:

```

train = all_data.loc[all_data['split']=='train']
train.head(5)

```

Out[]:

	id	comment_text	split	created_date	publication_id	parent_id	article_id	rating	funny	wow	sad	likes	disagree	toxicity	sev
0	1083994	He got his money... now he lies in wait till a...	train	2017-03-06 15:21:53.675241+00	21	NaN	317120	approved	0	0	0	2	0	0.373134	
1	650904	Mad dog will surely put the liberals in mental...	train	2016-12-02 16:44:21.329535+00	21	NaN	154086	approved	0	0	1	2	0	0.605263	
2	5902188	And Trump continues his lifelong cowardice by ...	train	2017-09-05 19:05:32.341360+00	55	NaN	374342	approved	1	0	2	3	7	0.666667	
4	5410943	Tucker and Paul are both total bad ass mofo's.	train	2017-06-14 05:08:21.997315+00	21	NaN	344096	approved	0	0	0	1	0	0.550000	
5	6290444	Cry me a river, why don't you.\nDrinking, drug...	train	2017-11-04 22:04:11.596185+00	54	6290143.0	396946	rejected	0	0	0	0	0	0.203390	



In []:

```
test = all_data.loc[all_data['split']=='test']
test.head(5)
```

Out[]:

	id	comment_text	split	created_date	publication_id	parent_id	article_id	rating	funny	wow	sad	likes	disagree	toxicity	se
3	7084460	"while arresting a man for resisting arrest".\...	test	2016-11-01 16:53:33.561631+00	13	NaN	149218	approved	0	0	0	0	0	0.815789	
10	7141509	NO ! There are no alternative facts. Go check...	test	2017-01-30 02:53:48.012277+00	21	919529.0	164687	approved	1	0	0	0	0	0.597222	
11	7077814	the more you whine sore loser Artster\n\nthe m...	test	2016-12-03 00:17:42.300700+00	54	649753.0	154126	approved	0	0	0	0	0	0.650000	
38	7147990	There's rarely opportunity to agree with Benne...	test	2017-09-13 16:37:16.990602+00	102	NaN	377304	approved	1	0	0	1	2	0.111111	
42	7008066	The Law has every freedom to be an ass!	test	2017-07-09 07:03:44.153492+00	54	5556167.0	353158	approved	0	0	0	0	0	0.800000	



In []:

```
train = train.reset_index(drop=True)
test = test.reset_index(drop=True)
```

In []:

```
X = train['clean_comment']
Y = train['sub_toxic']
```

In []:

```
x_train, x_test, y_train, y_test = train_test_split(X.values,Y.values, test_size=0.2, stratify=Y)
```

In []:

```
x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

Out[]:

```
((1443900,), (360975,), (1443900,), (360975,))
```

In []:

```
train_labels = to_categorical(y_train)
val_labels = to_categorical(y_test)
```

Loading BERT

In []:

```
bert_preprocess = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3")
bert_encoder = hub.KerasLayer("https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-512_A-8/2") #be:
```

In []:

```
bert = load_model('bert.hdf5', custom_objects={'KerasLayer': bert_preprocess})
```

In []:

```
bert.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
text (InputLayer)	[(None,)]	0	[]
keras_layer (KerasLayer)	{'input_type_ids': (None, 128), 'input_word_ids': (None, 128), 'input_mask': (None, 128)}	0	['text[0][0]']
keras_layer_1 (KerasLayer)	{'default': (None, 512), 'pooled_output': (None, 512), 'encoder_outputs': [(None, 128, 512), (None, 128, 512), (None, 128, 512), (None, 128, 512)], 'sequence_output': (None, 128, 512)}	28763649	['keras_layer[0][0]', 'keras_layer[0][1]', 'keras_layer[0][2]']
dropout (Dropout)	(None, 512)	0	['keras_layer_1[0][5]']
output (Dense)	(None, 7)	3591	['dropout[0][0]']
Total params: 28,767,240 Trainable params: 3,591 Non-trainable params: 28,763,649			

Train Prediction

In []:

```
train_prediction = bert.predict(train['clean_comment'])
```

In []:

```
train_classes = np.argmax(train_prediction,axis=1)
```

In []:

```
fone = f1_score(train['sub_toxic'].values, train_classes, average=None)
fone
```

Out []:


```
array([0.97108389, 0.0334355 , 0.15697937, 0.05373858, 0.01380814,
       0.03066038, 0.          ])
```

In []:

```

confusion = confusion_matrix(train['sub_toxic'].values, train_classes)
plt.figure(figsize = (16,10))
sns.heatmap(confusion, annot=True, fmt='g')
plt.xlabel('Actual values')
plt.ylabel('Predicted values')
plt.show()

```



Test Prediction

```

In []:
test_prediction = bert.predict(test['clean_comment'])


In []:
test_classes = np.argmax(test_prediction,axis=1)

In []:
fone2 = f1_score(test['sub_toxic'].values, test_classes, average=None)
fone2

Out[]:
array([0.9712429 , 0.00555556, 0.15759813, 0.04189636, 0.
        0.02659574, 0.
       ])

In []:
confusion = confusion_matrix(test['sub_toxic'].values, test_classes)
plt.figure(figsize = (16,10))
sns.heatmap(confusion, annot=True, fmt='g')
plt.xlabel('Actual values')
plt.ylabel('Predicted values')
plt.show()

```



Train: Incorrect predictions (Worst Case)

LIME analysis :Class 0

In []:

```
train['predicted'] = train_classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==0]
new_df = new_df[new_df['predicted'] != 0]
new_df.reset_index(drop=True, inplace=True)
```

In []:

```
idx = np.random.randint(0, len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 0
Predicted class: 2


In []:

```
class_names = ['non-toxic', 'threat', 'insult', 'identity_attack', 'sexual_explicit', 'obscene', 'severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10, labels=[0, 1, 2],
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities


insult	0.71
non-toxic	0.28
obscene	0.00
threat	0.00
Other	0.00

NOT non-toxic non-toxic




NOT threat

threat




NOT insult

insult




NOT identity_attack


identity_attack



NOT sexual_explicit

sexual_explicit





LIME analysis :Class 1

```
train['predicted'] = train_classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==1]
new_df = new_df[new_df['predicted'] != 1]
new_df.reset_index(drop=True, inplace=True)
```

In []:

```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```


In []:

```
Actual class: 1
Predicted class: 0
```


In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```


Prediction probabilities




NOT non-toxic non-toxic



NOT threat threat



NOT insult insult



NOT identity_attack identity_attack



men
care
n
every
t
gun
get
NOT sexual_explicit sexual_explicit

women
many
nut
due
men
get
n
every
care
killed

NOT obscene obscene

n
nut
every
get
guns
due
children
killed
keep
care

NOT severe_toxicity severe_toxicity

nut
n
killed
every
due
guns
children
t
keep
many

Text with highlighted words

n't care many men women children killed due guns long get keep guns every gun nut

LIME analysis :Class 2

In []:

```
train['predicted'] = train_classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==2]
new_df = new_df[new_df['predicted'] != 2]
new_df.reset_index(drop=True, inplace=True)
```

In []:


```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 2
Predicted class: 0


In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```


Prediction probabilities




NOT non-toxic non-toxic

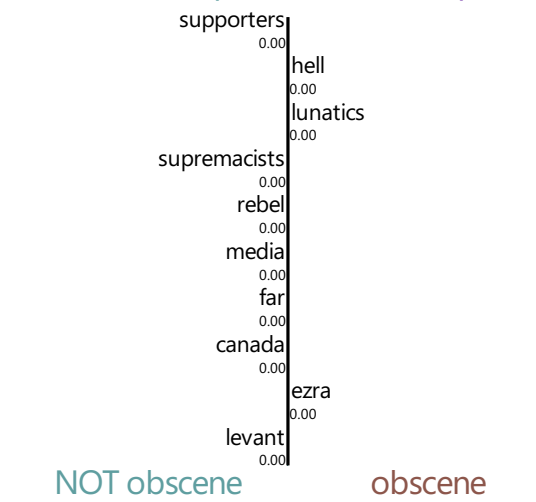
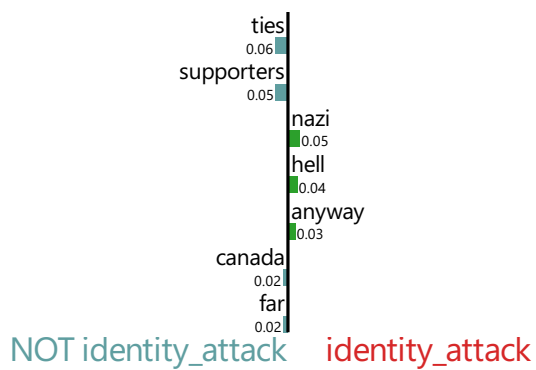



NOT threat threat



NOT insult insult







Text with highlighted words

rebel media ties right-wing groups n't impossible n't far right lunatics canada right right right anyway hell ezra levant sleep night knowing well far-right/alt-right white supremacists anti-semites nazi supporters ezra

LIME analysis :Class 3

In []:

```
train['predicted'] = train_classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==3]
new_df = new_df[new_df['predicted'] != 3]
new_df.reset_index(drop=True, inplace=True)
```

In []:


```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 3
Predicted class: 0

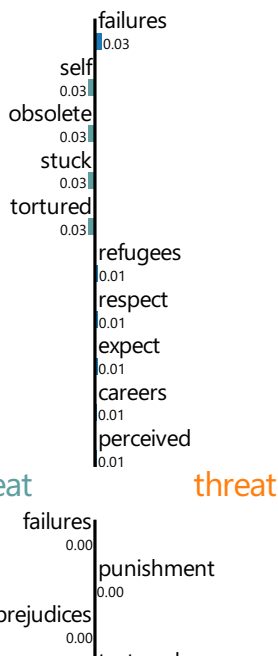
In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10, labels=[0, 1,2],
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities

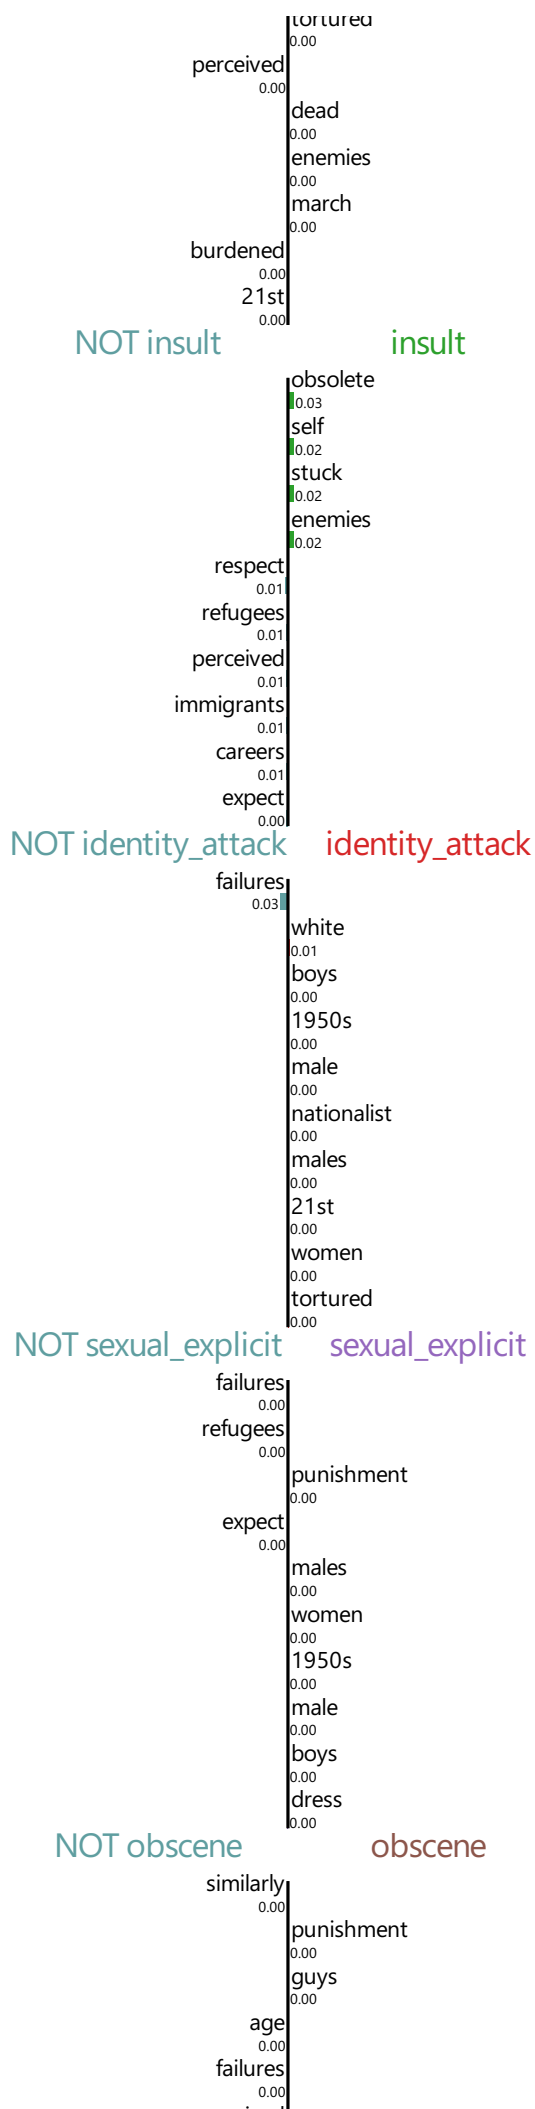


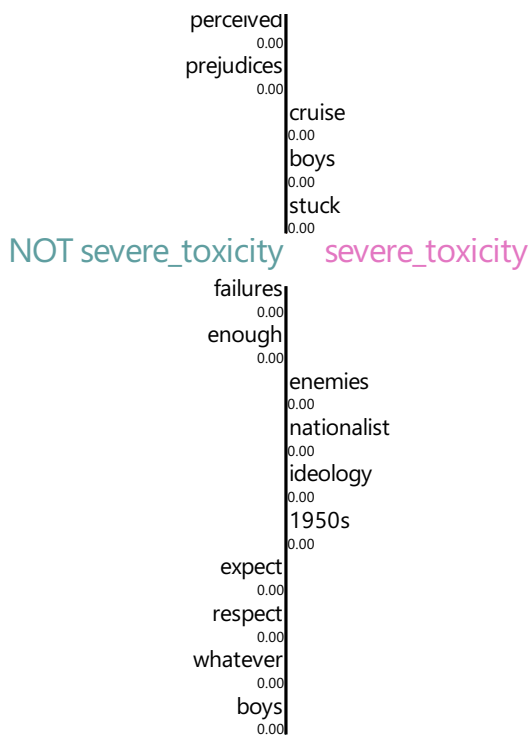
NOT non-toxic non-toxic



NOT threat

threat





Text with highlighted words

proud boys describe guys doofusses doofi embittered white males tortured loss privilege lashing perceived enemies refugees 1950s obsolete years age whatever classic fodder far right/white supremacist/alt right/nationalist fringe crushed simply male turned enough put front gravy train stuck dead-end no-hope careers watch women immigrants smart enough hard-working enough educate 21st century cruise past line dress prejudices expect taken seriously anyone similarly self-burdened brethren let march around little uniforms smaller ideology non-existent self-respect punishment failures

LIME analysis :Class 4

```
In [ ]:
train['predicted'] = train_classes

In [ ]:
new_df = train.loc[train['sub_toxic']==4]
new_df = new_df[new_df['predicted'] != 4]
new_df.reset_index(drop=True, inplace=True)

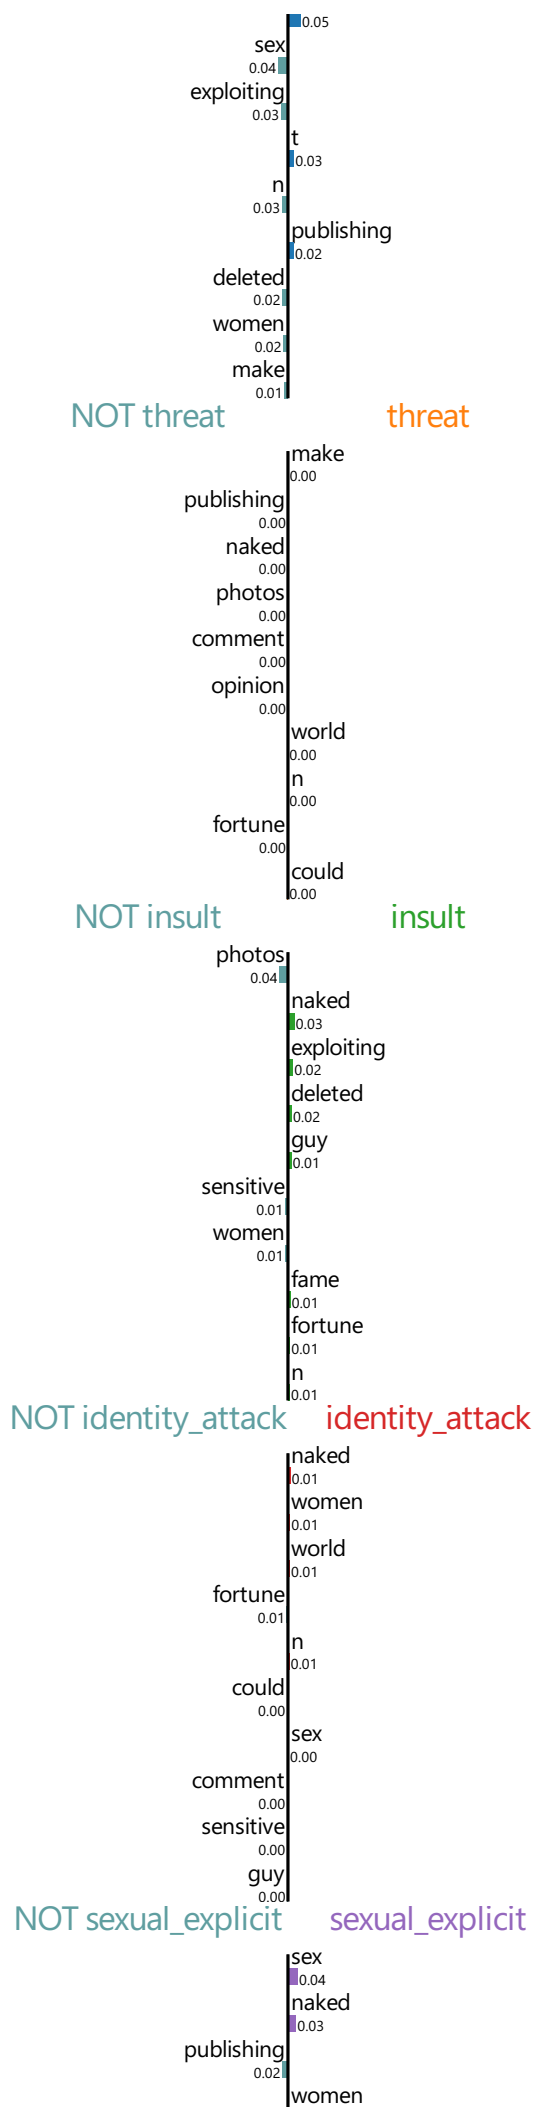
In [ ]:
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])

Actual class: 4
Predicted class: 0

In [ ]:
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities







Text with highlighted words

guy could make fortune fame sex publishing exploiting naked women photos n't make comment practice deleted opinion sensitive naked
 photos strange world

LIME analysis :Class 5

```
In []:
train['predicted'] = train_classes

In []:
new_df = train.loc[train['sub_toxic']==5]
new_df = new_df[new_df['predicted'] != 5]
new_df.reset_index(drop=True, inplace=True)


In []:
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])

Actual class:  5
Predicted class:  0


In []:
```

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```


Prediction probabilities




NOT non-toxic non-toxic




NOT threat threat

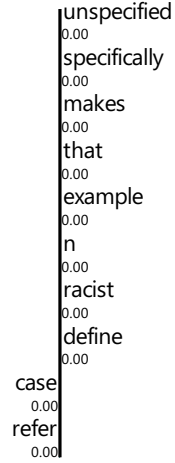
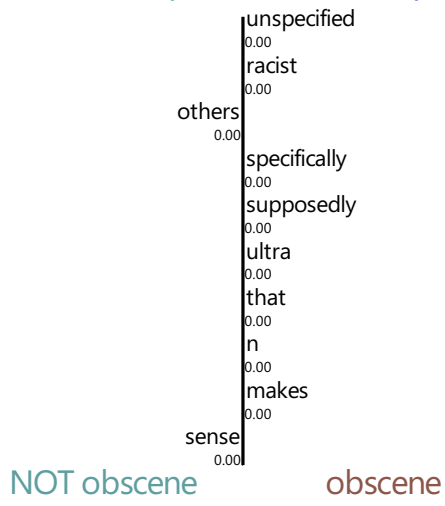



NOT insult insult



NOT identity_attack identity_attack





Text with highlighted words

unspecified 'crap' refer define 'specifically' makes whatever think 'that something comparable documented racism evident session 'case' reagan romney 'host others supposedly 'non-liberal 'specific follow politics pretty clearly done several decades enough try paint pictures big brush picture n't make sense unless add detail exactly say 'it said 'specifically' said example might john tepton 'specifically' makes refer 'ultra' racist

```
train.loc[train['sub_toxic']==6]
```

															Out[]:
	id	comment_text	split	created_date	publication_id	parent_id	article_id	rating	funny	wow	sad	likes	disagree	to	
7656	6214744	female assassins for kim	train	2017-10-24 19:47:26.537935+00	55	NaN	392890	approved	0	0	0	0	0	0	
537774	5801652	Awesome! Lets cut the head off hate! Lets stab...	train	2017-08-19 17:33:30.941935+00	21	NaN	368010	rejected	0	0	0	0	0	0	
1195376	741611	So dealers of death deserve clemency? What a w...	train	2016-12-22 07:13:30.077948+00	21	NaN	157754	approved	0	0	0	0	0	0	
1347411	5640242	I've worked in some estrogen swamps. Let me te...	train	2017-07-22 18:36:05.726092+00	54	NaN	357876	rejected	0	0	0	0	0	0	
1435818	5833628	Let me remind you people: If you disagree with...	train	2017-08-24 19:24:09.989776+00	54	NaN	370106	approved	3	0	0	0	0	0	
1450132	6214411	Peta is busy killing cats and dogs.\nhttps://w...	train	2017-10-24 19:08:54.717860+00	21	6212387.0	392564	approved	0	0	0	2	0	0	
1628400	5963721	Mr. Dries, Bertha Cooper represents the malign...	train	2017-09-16 12:21:43.299091+00	85	NaN	378495	rejected	0	0	0	0	0	0	

LIME analysis :Class 6

In []:

```
train['predicted'] = train_classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==6]
new_df = new_df[new_df['predicted'] != 6]
new_df.reset_index(drop=True, inplace=True)
```

In []:

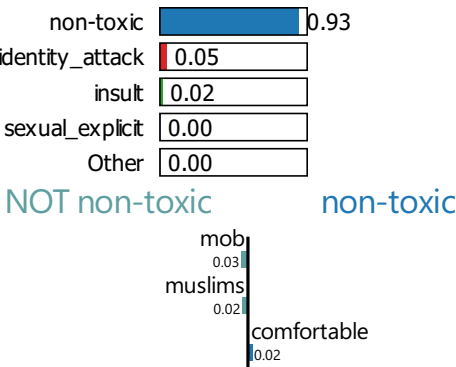
```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```


Actual class: 6
Predicted class: 0

In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities







Text with highlighted words

mr dries bertha cooper represents malignant mob never written anything designed undermine america promote muslims love move iran wait bombing know comfortable friends

Train : Correct Predictions (Strong Case)

LIME analysis :Class 0

```
train['predicted'] = train_classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==0]
new_df = new_df[new_df['predicted'] == 0]
new_df.reset_index(drop=True, inplace=True)
```

In []:

```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

In []:

Actual class: 0
Predicted class: 0


In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities

non-toxic	0.99
insult	0.01
obscene	0.00
sexual_explicit	0.00
Other	0.00


NOT non-toxic non-toxic




NOT threat threat



NOT insult insult



NOT identity_attack identity_attack



0.00
n
0.00
investor
0.00
suggested
0.00
property
0.00
sold
0.00
cheap
0.00
t
0.00
NOT sexual_explicit sexual_explicit

t
0.00
support
0.00
yes
0.00
source
0.00
n
0.00
capstone
0.00
profit
0.00
investor
0.00
corporation
0.00
cheap
0.00
NOT obscene obscene

suggested
0.00
support
0.00
investor
0.00
n
0.00
cheap
0.00
emption
0.00
sell
0.00
yes
0.00
profit
0.00
property
0.00
NOT severe_toxicity severe_toxicity

investment
0.00
t
0.00
n
0.00
yet
0.00
sold
0.00
yes
0.00
source
0.00
property
0.00
heard
0.00
emption
0.00

Text with highlighted words

yes heard credible source suggested capstone corporation n't sold housing project yet 're know build cheap property tax emption mupte
sell profit investor interested rents n't support investment

LIME analysis :Class 1

In []:

```
train['predicted'] = train_classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==1]
new_df = new_df[new_df['predicted'] == 1]
new_df.reset_index(drop=True, inplace=True)
```

In []:

```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

```
Actual class:  1
Predicted class:  1
```


In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```


Prediction probabilities

threat	0.54
non-toxic	0.45
identity_attack	0.01
insult	0.01
Other	0.00

NOT non-toxic non-toxic



NOT threat threat



NOT insult insult




NOT identity_attack identity_attack



NOT sexual_explicit sexual_explicit



NOT obscene obscene



NOT severe_toxicity severe_toxicity



Text with highlighted words

guns peoples homes shoot

LIME analysis :Class 2

In []:

```
train['predicted'] = train classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==2]
new_df = new_df[new_df['predicted'] == 2]
new_df.reset_index(drop=True, inplace=True)
```

In []:


```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

```
Actual class: 2
Predicted class: 2
```

In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10, labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```


Prediction probabilities



NOT non-toxic non-toxic



NOT threat threat



NOT insult insult



NOT identity_attack identity_attack



NOT sexual_explicit sexual_explicit



NOT obscene obscene



NOT severe_toxicity severe_toxicity



Text with highlighted words

fucking stupid canadian pig

LIME analysis :Class 3

In []:

```
train['predicted'] = train classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==3]
new_df = new_df[new_df['predicted'] == 3]
new_df.reset_index(drop=True, inplace=True)
```

In []:

```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

```
Actual class: 3
Predicted class: 3
```

In []:


```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities




NOT non-toxic

non-toxic



NOT threat

threat




NOT insult

insult




NOT identity_attack

identity_attack




NOT sexual_explicit

sexual_explicit




NOT obscene

obscene



NOT severe_toxicity

severe_toxicity



Text with highlighted words

muslims terrorists terrorists muslim

LIME analysis :Class 4

In []:

```
train['predicted'] = train_classes
```

In []:

```
new_df = train.loc[train['sub_toxic']==4]
new_df = new_df[new_df['predicted'] == 4]
new_df.reset_index(drop=True, inplace=True)
```

In []:


```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```


Actual class: 4
Predicted class: 4


In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities




NOT non-toxic non-toxic




NOT threat threat




NOT insult insult




NOT identity_attack identity_attack



NOT sexual_explicit sexual_explicit





Text with highlighted words

problem penis wanting use vagina restroom vagina wanting use penis restroom understand

LIME analysis :Class 5

```
In [ ]:
train['predicted'] = train_classes

In [ ]:
new_df = train.loc[train['sub_toxic']==5]
new_df = new_df[new_df['predicted'] == 5]
new_df.reset_index(drop=True, inplace=True)

In [ ]:
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])

Actual class:  5
Predicted class:  5

In [ ]:
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities



Text with highlighted words

fuck

Test: Incorrect predictions (Worst Case)

LIME analysis :Class 0

In []:

```
test['predicted'] = test_classes
```

In []:

```
new_df = test.loc[test['sub_toxic']==0]
new_df = new_df[new_df['predicted'] != 0]
new_df.reset_index(drop=True, inplace=True)
```

In []:

```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 0
Predicted class: 2

In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities

insult	0.65
non-toxic	0.33
obscene	0.02
threat	0.00
Other	0.00

NOT non-toxic non-toxic

gobbledygook 0.19
bunch 0.14

NOT threat threat

bunch 0.00
gobbledygook 0.00

NOT insult insult

gobbledygook 0.19
bunch 0.13

NOT identity_attack identity_attack

bunch 0.00
gobbledygook 0.00

NOT sexual_explicit sexual_explicit

gobbledygook 0.00
bunch 0.00

NOT obscene obscene

bunch 0.01
gobbledygook 0.00

NOT severe_toxicity severe_toxicity

bunch 0.00
gobbledygook 0.00

Text with highlighted words

bunch gobbledygook

LIME analysis :Class 1

In []:

```
test['predicted'] = test_classes
```

In []:

```
new_df = test.loc[test['sub_toxic']==1]
new_df = new_df[new_df['predicted'] != 1]
new_df.reset_index(drop=True, inplace=True)
```

In []:

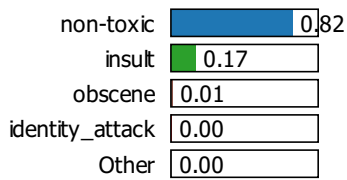
```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 1
Predicted class: 0


In []:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities




NOT non-toxic non-toxic




NOT threat threat


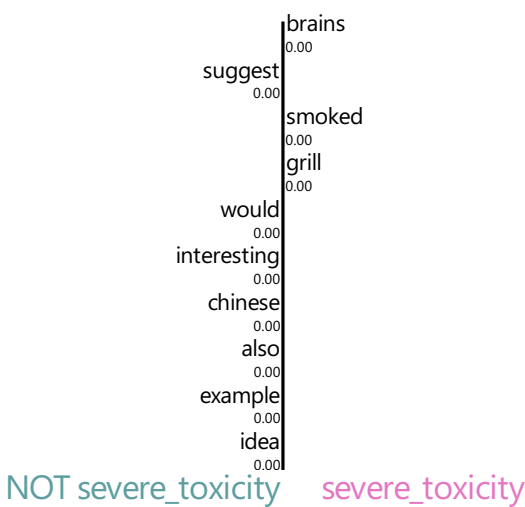


NOT insult insult



NOT identity_attack identity_attack





Text with highlighted words

chinese people **kill** cigarette smokers n't know torture like guantanamo james sayeth smith kline french n't portfolio pharmacy fred meyer patrons joes bar grill wholesale slaughter **also** mean perpetrating **evil** eyes really **expected** switch brands **would** suggest sereptishisley drank alcohol smoked ciggies one blow **brains** set good **example** **interesting** **idea** **mao**

LIME analysis :Class 2

```
In []:
test['predicted'] = test_classes

In []:
new_df = test.loc[test['sub_toxic']==2]
```

```
new_df = new_df[new_df['predicted'] != 2]
new_df.reset_index(drop=True, inplace=True)
```

In []:


```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

```
Actual class: 2
Predicted class: 0
```

In []:


```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities




NOT non-toxic

non-toxic




NOT threat

threat



NOT insult

insult



NOT identity_attack identity_attack

lots
0.02
knew
0.02
left
0.02
racist
0.04
today
0.02
left
0.02
misogynist
0.02
suicide
0.01
republicans
0.01
dim
0.01
republican
0.01
bigoted
0.01
party
0.01

NOT sexual_explicit sexual_explicit

republican
0.00
racist
0.00
misogynist
0.00
republicans
0.00
group
0.00
lots
0.00
watch
0.00
knew
0.00
left
0.00
today
0.00

NOT obscene obscene

knew
0.00
today
0.00
left
0.00
republican
0.00
bigoted
0.00
party
0.00
suicide
0.00
racist
0.00
republicans
0.00
bulbs
0.00

NOT severe_toxicity severe_toxicity

knew
0.00
lots
0.00
bigoted
0.00
left
0.00
racist
0.00
republican
0.00
suicide
0.00
republicans
0.00
group
0.00

party
0.00

Text with highlighted words

lots republicans suicide watch today republican party knew group racist misogynist bigoted dim bulbs left

LIME analysis :Class 3

In []:

```
test['predicted'] = test_classes
```

In []:

```
new_df = test.loc[test['sub_toxic']==3]
new_df = new_df[new_df['predicted'] != 3]
new_df.reset_index(drop=True, inplace=True)
```

In []:


```
idx = np.random.randint(0, len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 3
Predicted class: 0

In []:

```
class_names = ['non-toxic', 'threat', 'insult', 'identity_attack', 'sexual_explicit', 'obscene', 'severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10, labels=[0, 1, 2],
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities




NOT non-toxic

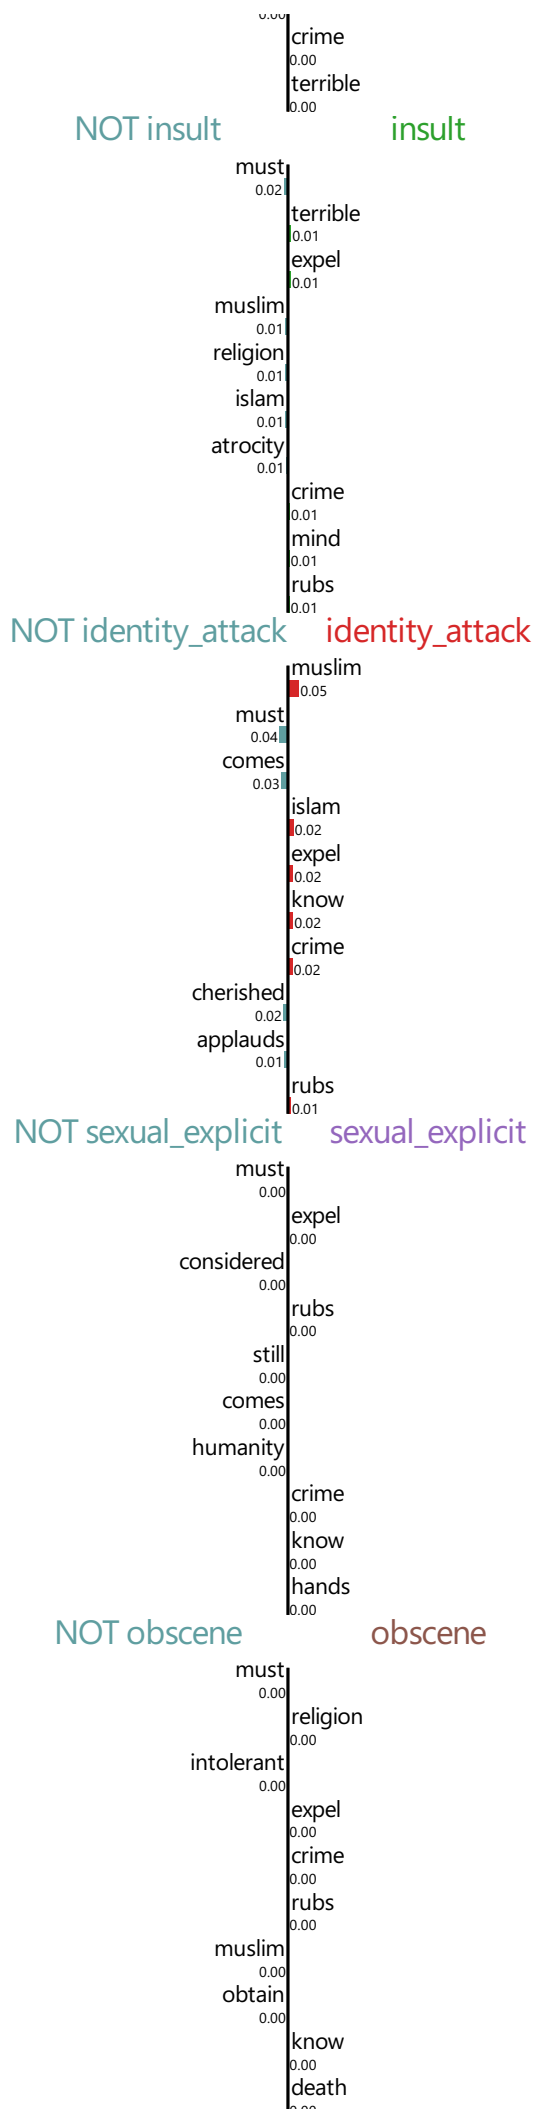
non-toxic



NOT threat

threat







Text with highlighted words

muslim causes terrible atrocity world still rubs hands islam religion considered one obtain religious freedom cherished west intolerant way life applauds death non believers world wide organized crime syndicate comes mind guise religion humanity must expel existence know

LIME analysis :Class 4

```
In []:
test['predicted'] = test_classes


In []:
new_df = test.loc[test['sub_toxic']==4]
new_df = new_df[new_df['predicted'] != 4]
new_df.reset_index(drop=True, inplace=True)

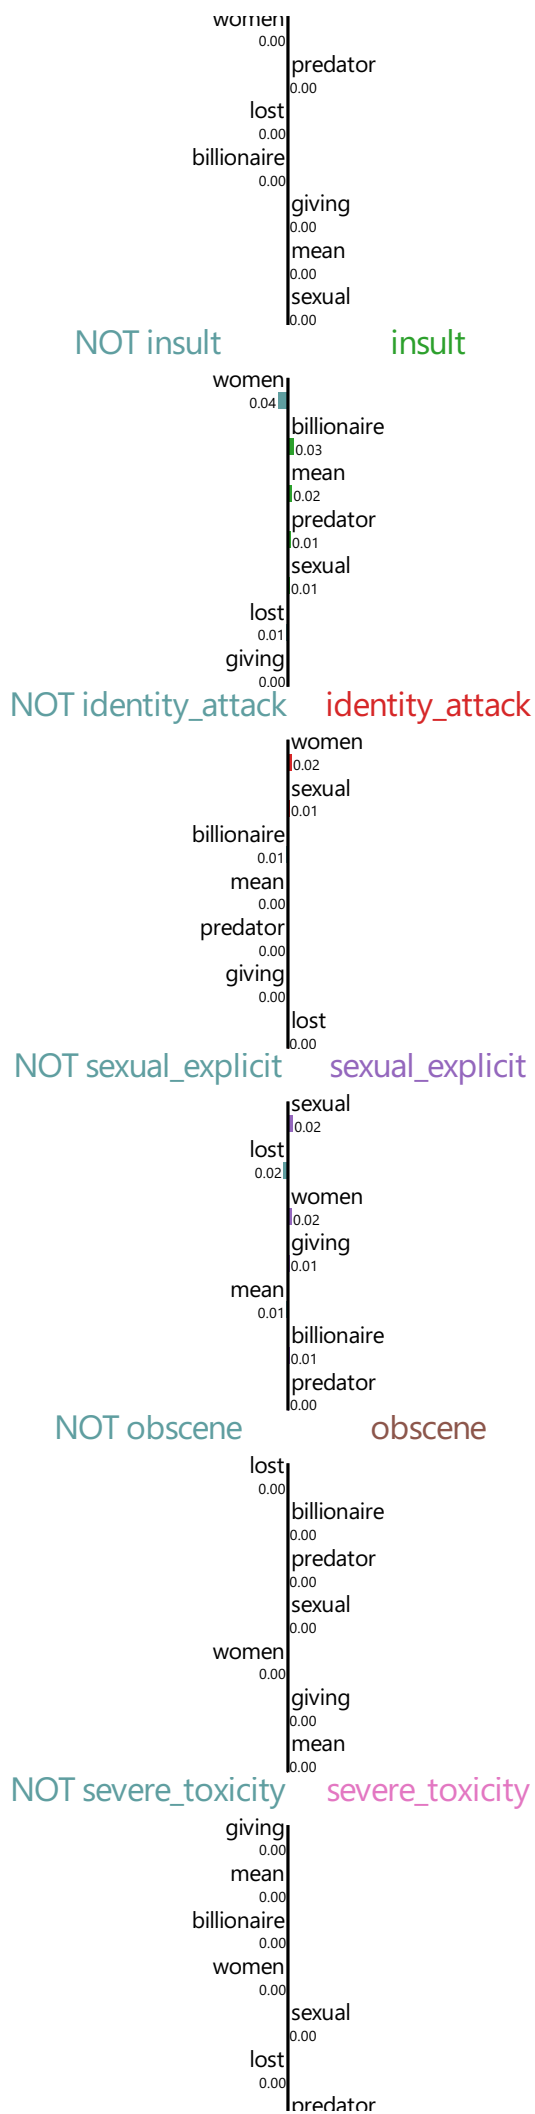
In []:
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])

Actual class: 4
Predicted class: 0

In []:
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities





Text with highlighted words

sexual predator mean women giving billionaire women lost

LIME analysis :Class 5

In []:

```
test['predicted'] = test_classes
```

In []:

```
new_df = test.loc[test['sub_toxic']==5]
new_df = new_df[new_df['predicted'] != 5]
new_df.reset_index(drop=True, inplace=True)
```

In []:

```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 5
Predicted class: 0

In []:

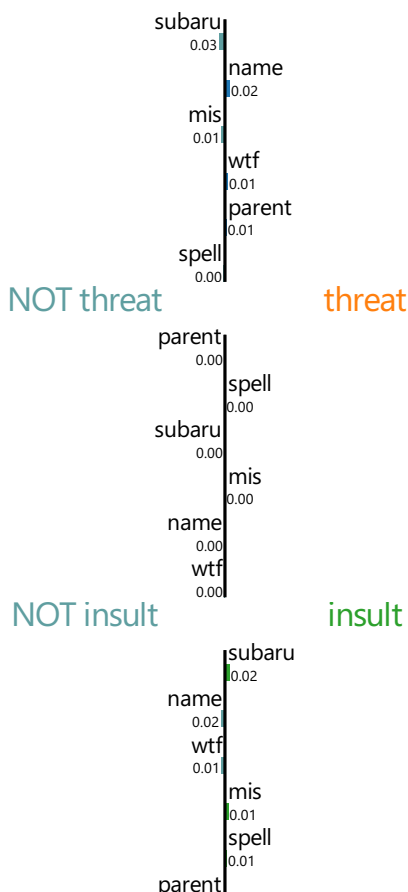
```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```


Prediction probabilities

non-toxic	0.96
insult	0.03
obscene	0.01
threat	0.00
Other	0.00

NOT non-toxic

non-toxic





Text with highlighted words

wtf name parent mis-spell subaru

LIME analysis :Class 6

In []:

```
test['predicted'] = test_classes
```

In []:

```
new_df = test.loc[test['sub_toxic']==6]
new_df = new_df[new_df['predicted'] != 6]
new_df.reset_index(drop=True, inplace=True)
```

In []:

```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```


Actual class: 6
Predicted class: 0

In []:


```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
```

```
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10, labels=[0, 1, 2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities




NOT non-toxic non-toxic




NOT threat threat




NOT insult insult



NOT identity_attack identity_attack





Text with highlighted words

n't insensitive young kids community died hell rail

Test : Correct Predictions (Strong Case)

LIME analysis :Class 0

```
In []:
test['predicted'] = test_classes

In []:

new_df = test.loc[test['sub_toxic']==0]
new_df = new_df[new_df['predicted'] == 0]
new_df.reset_index(drop=True, inplace=True)

In []:
```



```

idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])

```

```

Actual class:  0
Predicted class:  0

```

In []:

```

class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])


```

Prediction probabilities

non-toxic	<div><div></div></div>	0.99
insult	<div><div></div></div>	0.00
identity_attack	<div><div></div></div>	0.00
threat	<div><div></div></div>	0.00
Other	<div><div></div></div>	0.00


NOT non-toxic

non-toxic




NOT threat

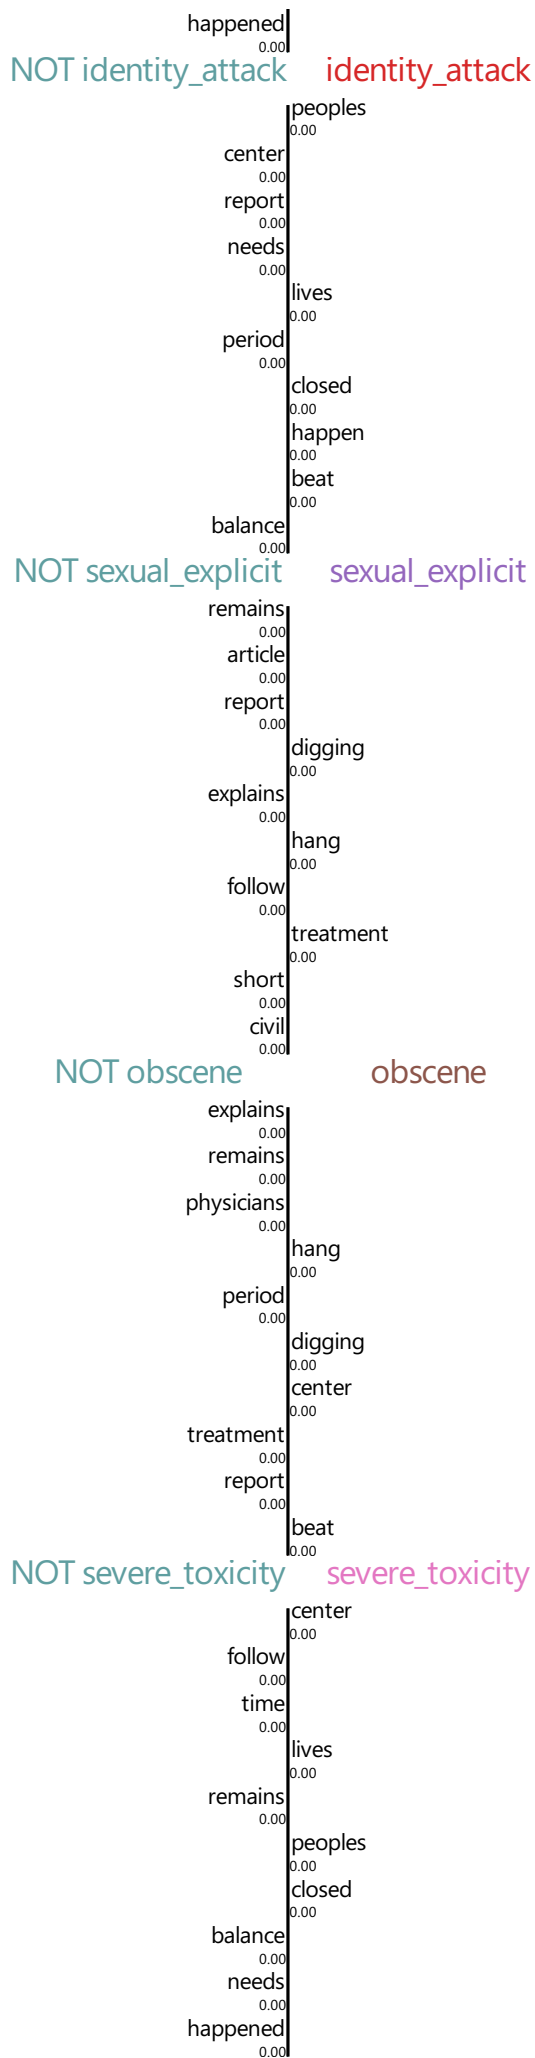
threat



NOT insult

insult





Text with highlighted words

article explains happened follow-up article needs report happen four physicians leave one another short period time come civil beat digging peoples lives hang balance long treatment center remains closed

LIME analysis :Class 1

In [198]:

```
test['predicted'] = test_classes
```

In [199]:

```
new_df = test.loc[test['sub_toxic']==1]
new_df = new_df[new_df['predicted'] == 1]
new_df.reset_index(drop=True, inplace=True)
```

In [200]:

```
new_df
```

Out[200]:

	id	comment_text	split	created_date	publication_id	parent_id	article_id	rating	funny	wow	sad	likes	disagree	toxicity	sev
0	7085319	If you pull a gun while committing a crime you...	test	2016-08-21 21:45:56.822437+00	21	NaN	144243	approved	0	0	0	7	0	0.581081	

In [201]:


```
idx= 0 #because only 1 index
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

```
Actual class: 1
Predicted class: 1
```


In [202]:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10, labels=[0, 1, 2],
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities



NOT non-toxic non-toxic



NOT threat threat





Text with highlighted words

pull gun committing crime deserve shot

LIME analysis :Class 2

In [194]:

```
test['predicted'] = test_classes
```

In [195]:

```
new_df = test.loc[test['sub_toxic']==2]
new_df = new_df[new_df['predicted'] == 2]
new_df.reset_index(drop=True, inplace=True)
```


In [196]:

```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 2
Predicted class: 2

In [197]:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```





Text with highlighted words

trump rest world everything persists stupidity

LIME analysis :Class 3

In [176]:

```
test['predicted'] = test_classes
```

In [177]:

```
new_df = test.loc[test['sub_toxic']==3]
new_df = new_df[new_df['predicted'] == 3]
new_df.reset_index(drop=True, inplace=True)
```

In [178]:


```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 3
Predicted class: 3


In [179]:

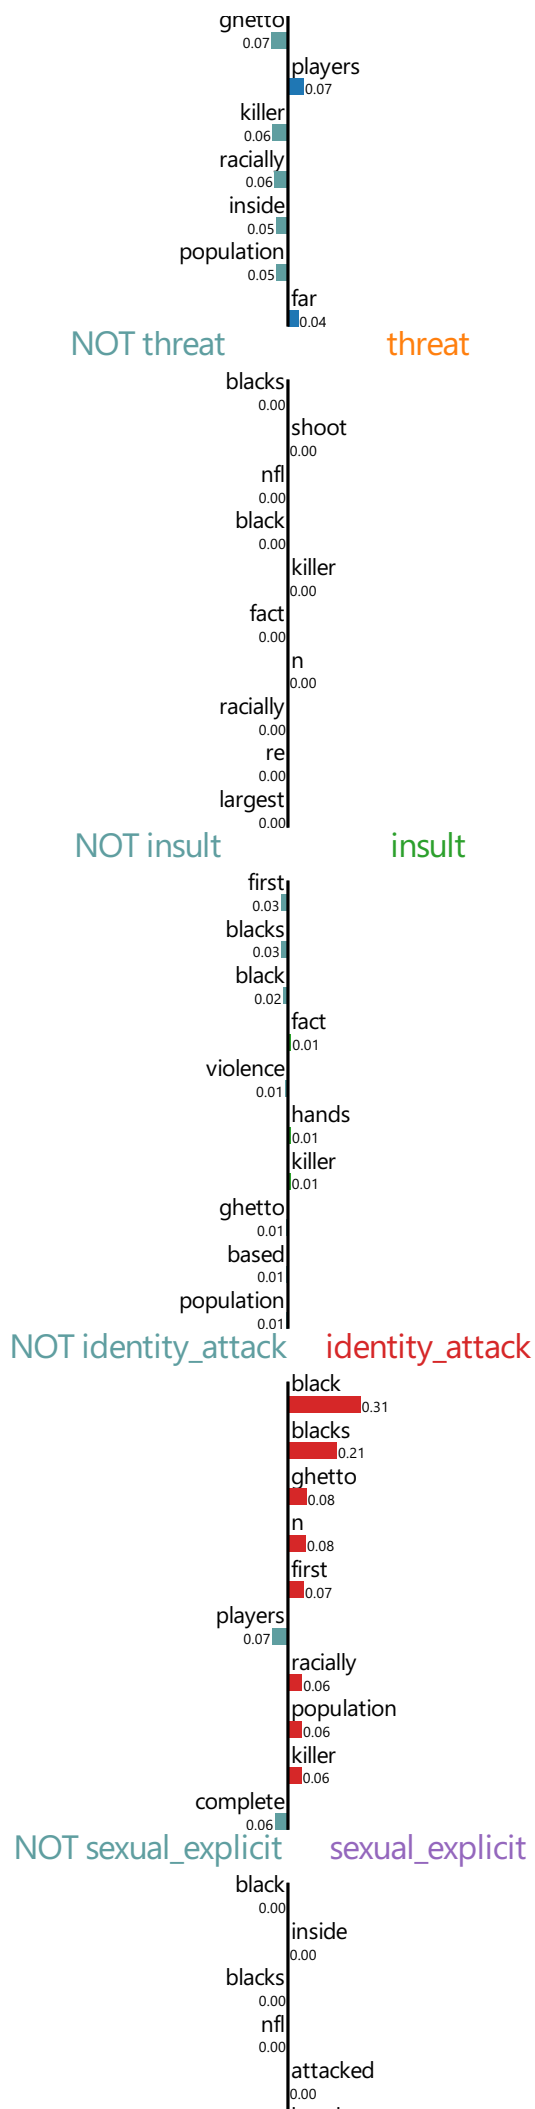
```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10,labels=[0, 1,2,
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```


Prediction probabilities



NOT non-toxic (left) and **non-toxic** (right) word importance charts.







Text with highlighted words

first 're protesting based lies blacks racially attacked law enforcement hands n't shoot complete lie nfl wants help black population need change ghetto culture inside black black violence far largest killer blacks nfl players n't touch fact

LIME analysis :Class 5

In [186]:

```
test['predicted'] = test_classes
```

In [187]:

```
new_df = test.loc[test['sub_toxic']==5]
new_df = new_df[new_df['predicted'] == 5]
new_df.reset_index(drop=True, inplace=True)
```

In [188]:

```
idx= np.random.random_integers(0,len(new_df))
print('Actual class: ', new_df['sub_toxic'][idx])
print('Predicted class: ', new_df['predicted'][idx])
```

Actual class: 5
Predicted class: 5


In [189]:

```
class_names = ['non-toxic','threat','insult','identity_attack','sexual_explicit','obscene','severe_toxicity']
explainer = LimeTextExplainer(class_names = class_names)
```




```
exp = explainer.explain_instance(new_df["clean_comment"][idx], bert.predict, num_features = 10, labels=[0, 1, 2],
exp.show_in_notebook(text=new_df["clean_comment"][idx])
```

Prediction probabilities




NOT non-toxic

non-toxic




NOT threat

threat




NOT insult

insult




NOT identity_attack

identity_attack




NOT sexual_explicit

sexual_explicit




NOT obscene

obscene



NOT severe_toxicity

severe_toxicity



Text with highlighted words

submarine damn

Points to be noted:-

- Class 0 non-toxic:-
 - Train - It is doing decent job in learning. Because of some words like murderous and traitor it incorrectly classified.
 - Test - It is doing good in classification. It did wrong classification because of some made up word.
- Class 1 threat:-
 - Test and Train - It is not doing that bad as the words because of which it is being classified as a threat are normal words and in the correct classification, we can see it is learning well to classify those words as threads.
- Class 2 insult:-
 - Train and Train - It is identifying well the words that belong to insult category but because of presence of more non-toxic words it is classifying them as non-toxic.
- Class 3 identity_attack:-
 - Train and Test - It is not giving weightage to identity words like islam and white. It is doing well to classify identity attack on blacks.
- Class 4 threat:-
 - Train and Test - It is not giving weightage to explicit words like naked, predator. It is doing well in knowing that they are toxic but due to presence of lots of non-toxic words it is classifying them as non-toxic.
- Class 5 obscene:-
 - Train and Test - It is doing well in identifying the obscene words like fuck but it is not classifying other obscene words like racist, crap well.
- Class 6 severe_toxicity:-
 - Train and Test - There are not enough examples for it to learn.

(Not included LIME analysis for any of the class in the sub-category of test or train where no data point was available)

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js